

Muhammad Safyan

Email: msafyano80@gmail.com

GitHub: github.com/MSafyan

LinkedIn: [linkedin.com/in/safyamakram](https://www.linkedin.com/in/safyamakram)

Cell: +923012772000

Location: Gujranwala, Pakistan

Portfolio: <https://safyan-portfolio.vercel.app>

SUMMARY

Full-stack and data engineer with 5+ years owning systems end-to-end: architecture, DevOps, and delivery. Shipped cloud-native backends, enterprise data warehouses, and AI/LLM platforms with measurable outcomes: 95% LLM cost reduction, 1M+ records/day ETL, \$40M portfolio visibility, 99.8%+ reliability.

EXPERIENCE

Pink Chicken (USA)

Mar 2025 – Jan 2026

Data Engineer | Azure · ADF · Shopify · Apparel Magic · Python · MsSQL · Dataflow

- Architected enterprise data warehouse on Azure integrating Shopify (10+ retail + e-commerce locations) with Apparel Magic WMS; designed virtual warehouse reconciliation logic to eliminate inventory double-counting
- Diagnosed 1–3% monthly sales divergence on multi-million-dollar revenue. Root cause: gift card refunds classified inconsistently across Shopify REST vs GraphQL; built reconciliation pipeline using GraphQL order changes API with explicit transaction classification logic; fixed discrepancy permanently
- Discovered incoming stock only exposed in Shopify GraphQL (not REST); redesigned inventory pipeline to merge both endpoints, resolving schema drift without disrupting existing incremental loads
- Reduced Azure Monitor costs by 99.99%, from ~\$1,000/month to \$0.10/month (~\$12,000/year saved) by eliminating redundant alert rules and consolidating to a single essential pipeline failure alert
- Implemented daily incremental loads, daily inventory snapshots, and dimensional model (fact/dimension/bridge tables) for multi-perspective inventory valuation across finance, D2C, and wholesale views

Palmetto Air Balance (USA)

Dec 2024 – Jan 2026

Data Engineer | Azure · ADF · NetSuite · UKG · Air1 · T-SQL

- Architected Azure-based enterprise warehouse: 50+ ADF pipelines integrating NetSuite (ERP), Air1 (project management), and UKG (workforce) with complex cross-system key mapping and financial reconciliation
- Delivered real-time visibility into \$40M+ project portfolio (revenue, backlog, utilization dashboards); reduced manual reporting effort by 70% across the business
- Built dynamic MERGE stored procedures with intelligent primary key detection, TRY-CATCH error handling, pipeline-level + database-level logging, email alerts with failure categorization, and audit trails
- Processed 1M+ records daily with sub-second query latency using composite indexing, table partitioning, and materialized views; implemented weekly automated backlog snapshots for historical trending

Abundance Movement (UK)

Apr 2023 – Present

Full-Stack + AI Engineer | NestJS · OpenAI · LangChain · Pinecone · TypeORM · PostgreSQL · AWS

- Owned full development lifecycle: designed and built NestJS backend from scratch with TypeORM + PostgreSQL, schema design, location-based indexes, ATM contract integration, and AWS deployment; Python Lambda for automated image/video compression on S3 with event triggers
- Architected 11-service AI recommendation and matching platform using OpenAI GPT, LangChain, and Pinecone; multi-index namespace architecture storing 5+ profile dimensions per user (strengths, weaknesses, interests, values, keywords) enabling semantic similarity matching at 35%+ threshold
- Reduced LLM costs by 95% through strategic model switching (o1-mini → gpt-4o-mini), token limiting (300–2,000 tokens per use case), and temperature optimization; maintained output quality throughout
- Built automated chat summarization pipeline: NestJS cron collects 7-day message windows, generates narrative summaries + action items via OpenAI, with REST endpoints for on-demand and historical digests

Helpicon (SWEDEN)

Sep 2023 – Mar 2026

Senior Backend + DevOps Engineer | Node.js · AWS ECS/EKS · Terraform · Jenkins · Docker

- Integrated Stripe and Swish (Nordic mobile payments) marketplace flows: automated provider commission splits, subscription tiers, and webhook handlers (payment_intent, subscriptions, refunds); achieved 99.8% payment success rate across 10,000+ active users and 500+ verified service providers
- Cut Twilio messaging costs by 90% via Conversations API shared-service architecture; improved API response times 2–3x to under 200ms via selective field projection, caching, and composite indexes; DB queries <50ms
- ECS auto-scaling (min 2 → max 10 at 70% CPU); Multi-AZ RDS PostgreSQL; RDS backup and snapshots;
- Deployed payment microservice on AWS EKS with Jenkins CI/CD across 3 isolated AWS accounts (DevOps / Staging / Production); Terraform IaC for RDS, ECR, S3, VPC; automated Prisma migrations

CVR Sports Analytics (AUS)

2022 – 2026

Full-Stack + Data Engineer | Node.js · MongoDB · Angular 17 · Express

- Designed analytics platform for NRL, Super League, and NSW Cup with a CQRS-like architecture: materialized read models at player/team/quarter/10-min grains. Event-derived KPIs computed once at XML ingestion, eliminating repeated recalculation on every report query
- MongoDB: compound indexes on real access patterns, parallel Promise.all for independent reads, distinct() for filter endpoints, lean() on non-hydration paths, aggregation pipelines close to data, self-join via \$lookup for team-vs-opponent metrics, allowDiskUse + maxTimeMS guardrails on heavy pipelines
- Led Angular 9 → 17 upgrade: removed legacy OpenSSL dependency, pruned unused packages, updated AG Grid with scoped theme overrides to fix visual regressions
- Engineered PDF export (html2canvas + jsPDF) across all report views; identified DOM rasterization as dominant cost; used onclone mutation to inject page-break spacers without touching the live page

E-Commerce Intelligence Platform (USA)

2025

Data Engineer | Snowflake · Snowpark Python · Keepa API · ScaleSERP

- Designed idempotent ETL on Snowflake + Snowpark processing 50k–100k+ products per run; achieved 75%+ catalog coverage vs 15% with single-strategy ingestion by implementing 6 configurable sort strategies around Keepa's API pagination constraints
- Batch INSERT...UNION ALL patterns eliminated per-row MAX(ID) lookups, achieving 10–20x performance improvement; canonical ENTITY/ATTRIBUTE model for semi-structured payloads; token-based dynamic rate limiting

Frizhub (PAK)

Nov 2022 – Aug 2023

Full-Stack Developer | Angular · Next.js · Web3.js · PostgreSQL · Alchemy API

- Built CHAINSCAN Ethereum NFT analytics platform: migrated millions of blockchain records from Alchemy Graph API into PostgreSQL with composite index optimization; implemented Web3 auth and wallet connectivity

P2G Global Services (USA)

Nov 2021 – 2023

- Automated management of 130k+ product catalog; Python scripts for inventory updates, order shipment tracking (Didadi APIs), SEO features, and Google Analytics segmentation

CERTIFICATIONS

AWS Certified Solutions Architect – Associate (2025) [Verified via Credly](#)
HashiCorp Certified: Terraform Associate (2024) [Verified via Credly](#)
Certified Kubernetes Application Developer – CKAD (2026) [Verified via Credly](#)

TECHNICAL SKILLS

Languages:	TypeScript, JavaScript, Python, T-SQL
Backend:	Node.js, NestJS, Express, REST APIs, Microservices, GraphQL
Frontend:	React.js, Next.js, Angular, TanStack Query, Redux, HTML5, CSS3/Sass
Databases:	PostgreSQL, Azure SQL, MongoDB, Snowflake, Redis, Pinecone (vector DB)
Cloud & DevOps:	AWS (EKS, ECS, RDS, S3, Lambda, CloudWatch), Azure (ADF, SQL, Storage), Terraform, Docker, Kubernetes,, GitHub Actions, Bitbucket Pipelines
AI/ML:	OpenAI GPT, LangChain, Pinecone, text embeddings, prompt engineering, cost optimization
Data & ETL:	Azure Data Factory, Snowpark, dimensional modeling
Commons:	Stripe, Swish, PayPal, Twilio Conversations, SendGrid

EDUCATION

BS Software Engineering, COMSATS University Islamabad, Lahore Campus